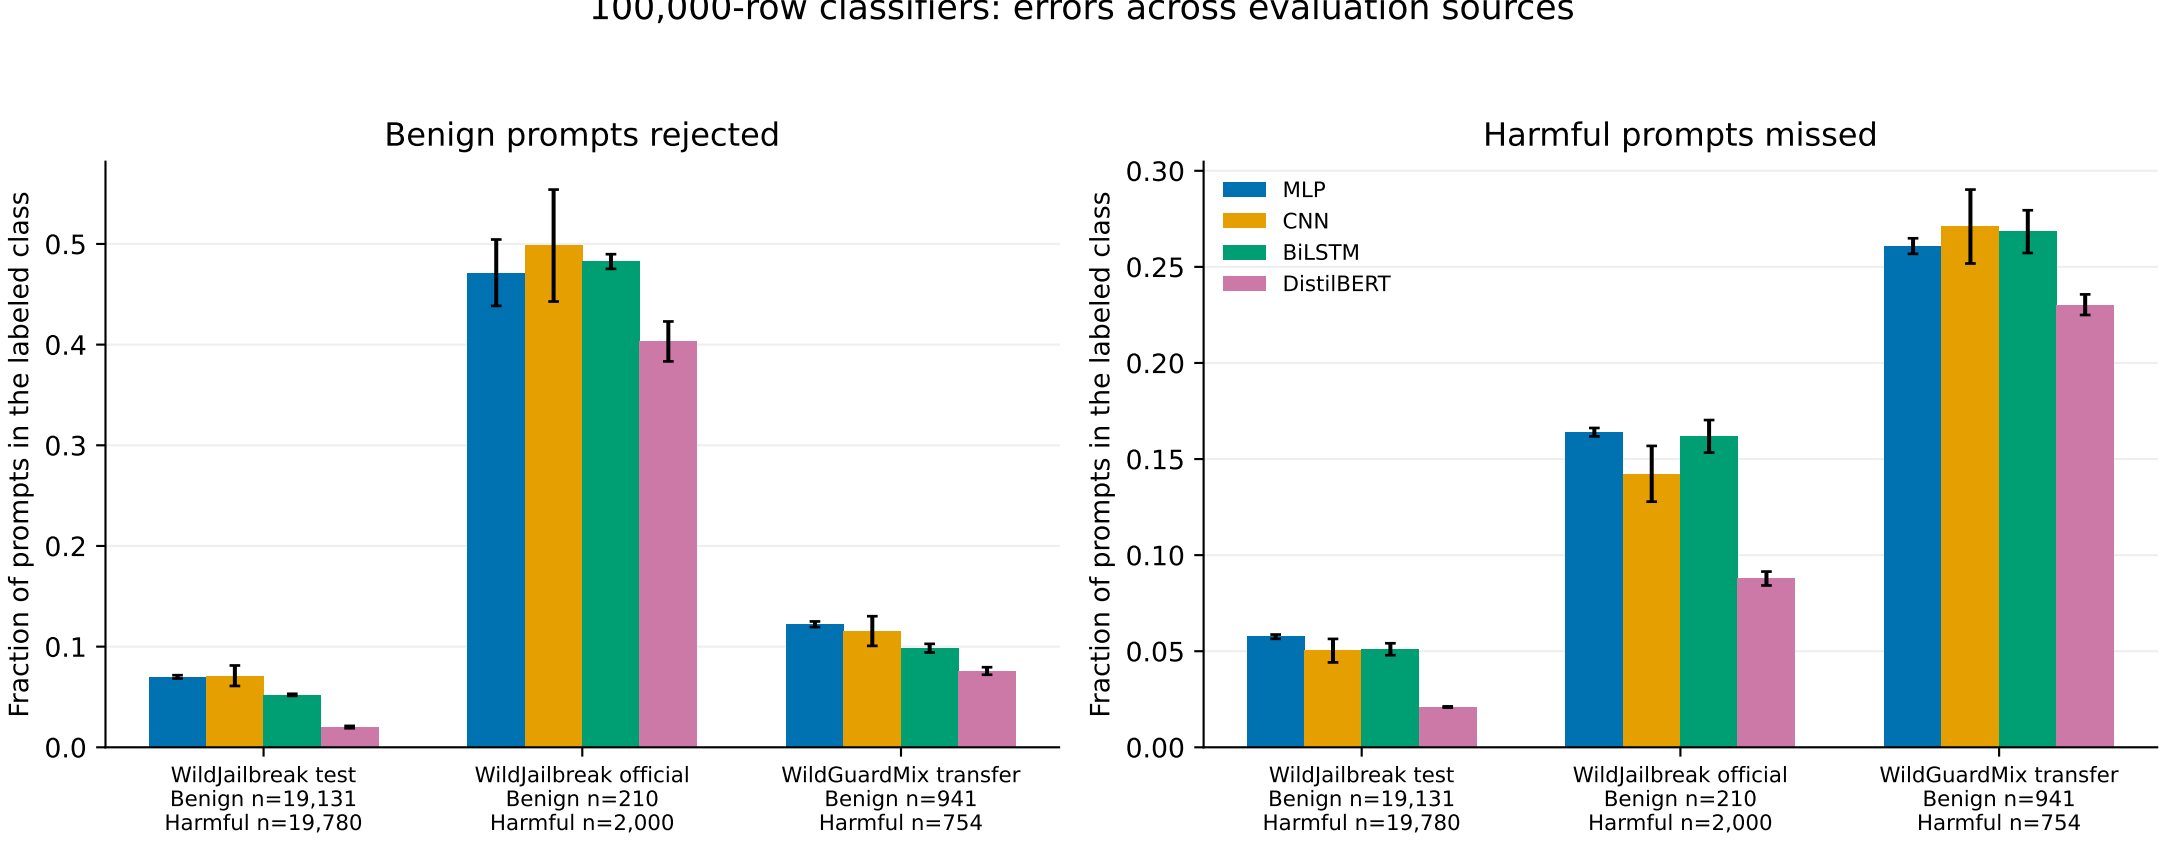
100,000-row classifiers: errors across evaluation sources
Benign prompts rejected
Harmful prompts missed
Fraction of prompts in the labeled class
MLP
CNN
BiLSTM
DistilBERT
0.0
0.1
0.2
0.3
0.4
0.5
0.00
0.05
0.10
0.15
0.20
0.25
0.30
WildJailbreak test
Benign n=19,131
Harmful n=19,780
WildJailbreak official
Benign n=210
Harmful n=2,000
WildGuardMix transfer
Benign n=941
Harmful n=754
Threshold 0.5; means and sample standard deviations across seeds 4442–4444. Evaluation distributions differ.