



# Prompt harmfulness detection · 100,000 unique training prompts

## Main grouped test: 38,911 prompts



## Official challenge: 2,000 harmful / 210 benign



Three epochs; validation-selected checkpoints; one training seed. F1 intervals: 95% bootstrap over request groups.